

Unit 4: Identifying Trends and Creating Models

Contents

- [Getting Started](#)
- [Linear Models](#)
 - [An Example](#)
 - [Simple Linear Regression](#)
 - [Multiple Linear Regression](#)
 - [Linear-Like Regression](#)
- [Logistic Regression](#)
 - [An Example](#)
 - [Home Data](#)
- [Lab Answers](#)
- [Next Steps](#)
- [Resources and Further Reading](#)
- [Notes](#)
- [Exercises](#)

Lab Questions

[1](#), [2](#), [3](#), [4](#), [5](#), [6](#), [7](#), [8](#)

Getting Started

In previous units, we worked on loading, cleaning, and exploring data. While working with the data, we noted that certain relationships appeared to exist between columns/variables. While plots allowed us to make claims like "x increases as y decreases", we didn't try to model that relationship mathematically nor did we try to determine the quality of that model.

In this unit, we'll look at creating models for the relationships in our data; specifically, we'll look at linear models for numerical data and logistic models for categorical data.

To create and explore these models, we'll use the [StatsModels](https://www.statsmodels.org/stable/index.html) (<https://www.statsmodels.org/stable/index.html>) library. To install it, we'll use `!pip .`

```
In [101]: import sys
!{sys.executable} -m pip install statsmodels
```

```
Requirement already satisfied: statsmodels in /usr/local/lib/python3.6/site-packages
Requirement already satisfied: pandas in /usr/local/lib/python3.6/site-packages (from statsmodels)
Requirement already satisfied: patsy in /usr/local/lib/python3.6/site-packages (from statsmodels)
Requirement already satisfied: scipy in /usr/local/lib/python3.6/site-packages (from statsmodels)
Requirement already satisfied: python-dateutil>=2 in /usr/local/lib/python3.6/site-packages (from pandas->statsmodels)
Requirement already satisfied: numpy>=1.9.0 in /usr/local/lib/python3.6/site-packages (from pandas->statsmodels)
Requirement already satisfied: pytz>=2011k in /usr/local/lib/python3.6/site-packages (from pandas->statsmodels)
Requirement already satisfied: six in /usr/local/lib/python3.6/site-packages (from patsy->statsmodels)
You are using pip version 9.0.3, however version 10.0.0 is available.
You should consider upgrading via the 'pip install --upgrade pip' command.
```

We'll work with plots in this unit. To ensure they are displayed in the notebook itself, we'll need to use the `%matplotlib inline` command.

```
In [102]: %matplotlib inline
```

In addition to the Seaborn and pandas libraries, we'll make explicit use of libraries on which they depend.

- [matplotlib.pyplot](https://matplotlib.org/api/pyplot_api.html) (https://matplotlib.org/api/pyplot_api.html): a collection of plotting functions with [MATLAB](https://www.mathworks.com/products/matlab.html) (<https://www.mathworks.com/products/matlab.html>)-like syntax
- [numpy](http://www.numpy.org/) (<http://www.numpy.org/>): scientific computing library

We'll import these and StatsModels with names that follow standard convention.

We also set the figure size for plots and a marker size for outliers in box plots.

```
In [103]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import statsmodels.api as sm

sns.set(rc={'figure.figsize':(12,8), "lines.markeredgewidth": 0.5 })
```

Linear Models

The first type of model we'll look at are linear models also known as linear regressions. In basic

terms, we use a linear model if the data looks like a line could be drawn through it. More specifically, for two dimensional data, we try to model the relationship between two variables, the independent or input variable, X , and the dependent or response variable, Y , by an equation of the form

$$Y = \beta_0 + \beta_1 X$$

where β_0 and β_1 are *coefficients*. We can generalize this to more than two dimensions. If X_1, X_2, \dots, X_n are independent variables, and Y is the dependent variable, we try to model the relationship by an equation in the form

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_n X_n$$

where $\beta_0, \beta_1, \dots, \beta_n$ are the coefficients.

To find the coefficients of these equations, we'll rely on the [ordinary least squares](https://en.wikipedia.org/wiki/Ordinary_least_squares) (https://en.wikipedia.org/wiki/Ordinary_least_squares) method that attempts to minimizing the the sum of squares of the differences between the observed values and the predicted values - we'll explore this further in a bit.

An alternative formulation of this problem involves vectors and matrices. For each observation, that is for each pair of values of $(x_0, y_0), (x_1, y_1), \dots, (x_m, y_m)$ we have the following system of equations:

$$\begin{aligned} y_0 &= \beta_0 + \beta_1 x_0 \\ y_1 &= \beta_0 + \beta_1 x_1 \\ &\vdots \\ y_m &= \beta_0 + \beta_1 x_m \end{aligned}$$

We can write this in terms of vectors.

$$\begin{bmatrix} y_0 \\ y_1 \\ \vdots \\ y_m \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix} \beta_0 + \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_m \end{bmatrix} \beta_1$$

We can write this in terms of matrix multiplication. We also swap the left and right hand sides.

$$\begin{bmatrix} 1 & x_0 \\ 1 & x_1 \\ \vdots & \vdots \\ 1 & x_m \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} = \begin{bmatrix} y_0 \\ y_1 \\ \vdots \\ y_m \end{bmatrix}$$

This is typically an overdetermined linear system of equations of the form

$$\mathbf{X}\beta = \mathbf{y}$$

where \mathbf{X} is the matrix of ones and observed values of the independent variable, β is the vector of unknown coefficients, and \mathbf{y} is the vector of observed values of the dependent variable. A variety of methods exist to find meaningful solutions to this linear equation. While this formulation is equivalent to ordinary least squares, it is commonly referred to as [linear least squares](https://en.wikipedia.org/wiki/Linear_least_squares_(mathematics)#Computation) ([https://en.wikipedia.org/wiki/Linear_least_squares_\(mathematics\)#Computation](https://en.wikipedia.org/wiki/Linear_least_squares_(mathematics)#Computation)).

An Example

Let's look at an example of how we can calculate the coefficients of a linear equation that models some data. We'll start with an example based on the [StatsModels documentation](http://www.statsmodels.org/stable/examples/notebooks/generated/ols.html) (<http://www.statsmodels.org/stable/examples/notebooks/generated/ols.html>).

First, we generate our "observed" data. To do this, we'll use the Numpy [linspace\(\)](https://docs.scipy.org/doc/numpy/reference/generated/numpy.linspace.html) (<https://docs.scipy.org/doc/numpy/reference/generated/numpy.linspace.html>) function to generate 100 evenly-spaced values between 0 and 10; these will correspond to values that will be used for the independent variable. Next, we'll use the [normal\(\)](https://docs.scipy.org/doc/numpy/reference/generated/numpy.random.normal.html) (<https://docs.scipy.org/doc/numpy/reference/generated/numpy.random.normal.html>) function from NumPy's [random](https://docs.scipy.org/doc/numpy-1.14.0/reference/routines.random.html) (<https://docs.scipy.org/doc/numpy-1.14.0/reference/routines.random.html>) submodule to draw 100 samples from a normal distribution - this will simulate errors in our data.

```
In [104]: nsample = 100
x = np.linspace(0, 10, nsample)
e = np.random.normal(size=nsample)

display(x[:10])

array([0.          , 0.1010101 , 0.2020202 , 0.3030303 , 0.4040404 ,
       0.50505051, 0.60606061, 0.70707071, 0.80808081, 0.90909091])
```

Looking at the first 10 values of `x`, we can see that they are stored in an [array](https://docs.scipy.org/doc/numpy-1.14.0/reference/generated/numpy.array.html). (<https://docs.scipy.org/doc/numpy-1.14.0/reference/generated/numpy.array.html>).

Next, we calculate our "observed" values as a combination of the independent variable, a constant, and some error.


```
In [105]: y = 5 * x + 3 + e
```

Let's plot the values of `x` and `y`. We'll create a scatter plot but use a different approach to do this.

First, we create [figure](https://matplotlib.org/api/_as_gen/matplotlib.figure.Figure.html#matplotlib.figure.Figure) (https://matplotlib.org/api/_as_gen/matplotlib.figure.Figure.html#matplotlib.figure.Figure) and [axes](https://matplotlib.org/api/axes_api.html#matplotlib.axes.Axes) (https://matplotlib.org/api/axes_api.html#matplotlib.axes.Axes) objects using the pyplot [subplots\(\)](https://matplotlib.org/api/_as_gen/matplotlib.pyplot.subplots.html) (https://matplotlib.org/api/_as_gen/matplotlib.pyplot.subplots.html) function; this is useful when we want to plot items from different sources together (as we will do in a bit). We use the axes' [plot\(\)](https://matplotlib.org/api/_as_gen/matplotlib.axes.Axes.plot.html#matplotlib.axes.Axes.plot) (https://matplotlib.org/api/_as_gen/matplotlib.axes.Axes.plot.html#matplotlib.axes.Axes.plot) method to plot the coordinate pairs from `x` and `y`; the third argument indicates that we'd like to use blue plus-sign markers rather than draw lines from the data.

```
In [106]: fig, ax = plt.subplots()
          ax.plot(x, y, 'b+')
```

```
Out[106]: [<matplotlib.lines.Line2D at 0x1129b1cf8>]
```



As shown by the plot, the data does look linearly related. One measure of the strength of the relationship is the [correlation coefficient](https://en.wikipedia.org/wiki/Correlation_coefficient) (https://en.wikipedia.org/wiki/Correlation_coefficient). Given two variables, the correlation coefficient can have a value between -1 and 1 where -1 indicates strong, negative relationship (as one variable increases, the other decreases) and 1 indicates a strong, positive relationship (as one variable increases, the other increases); a value of zero indicates no relationship exists.

To calculate the correlation coefficient, we can use the NumPy [corrcoef\(\)](https://docs.scipy.org/doc/numpy/reference/generated/numpy.corrcoef.html) (<https://docs.scipy.org/doc/numpy/reference/generated/numpy.corrcoef.html>) function.

```
In [107]: np.corrcoef(x, y)
```

```
Out[107]: array([[1.          , 0.99743872],
                 [0.99743872, 1.          ]])
```

The output is the correlation coefficient matrix that shows from left to right, top to bottom, correlation coefficient between the first variable and itself, the correlation coefficient between the first variable and the second, the correlation coefficient between the second variable and the first, and the correlation coefficient between the second variable and itself. We expect that a variable will be strongly correlated with itself. The output indicates a strong relationship between the variables. By construction, the relationship is not only strong, it is also very linear. To see this, we can use a linear regression.

In the two-dimensional linear regression, we need to calculate the values of two coefficients: a constant and a value that will be multiplied by the value of the independent variable. As we saw above, we can write the linear model in the form

$$Y = \beta_0 + \beta_1 X$$

We can write the constant, β_0 as $\beta_0 X^0$. Since any value raised to the zeroth power is one, we can write β_0 as $1 \cdot \beta_0$. We can say that the dependent variable is a linear combination of 1 and the independent variable. For our model, we account for this by using the StatsModels' [add_constant\(\)](http://www.statsmodels.org/dev/generated/statsmodels.tools.tools.add_constant.html) (http://www.statsmodels.org/dev/generated/statsmodels.tools.tools.add_constant.html) function.

```
In [108]: X = sm.add_constant(x)
display(X[:10])

array([[1.          , 0.          ],
       [1.          , 0.1010101 ],
       [1.          , 0.2020202 ],
       [1.          , 0.3030303 ],
       [1.          , 0.4040404 ],
       [1.          , 0.5050505 ],
       [1.          , 0.6060606 ],
       [1.          , 0.7070707 ],
       [1.          , 0.8080808 ],
       [1.          , 0.9090909 ]])
```

Comparing `x` to `X` we now have an array of arrays where the first element in the inner arrays are 1 corresponding to the value of the independent variable for the constant term. The corresponds to the matrix **X** described above.

To compute the the linear regression, we first set up the ordinary least squares model using the StatsModels' [OLS](http://www.statsmodels.org/dev/generated/statsmodels.regression.linear_model.OLS.html) (http://www.statsmodels.org/dev/generated/statsmodels.regression.linear_model.OLS.html) function and by specifying the array of values for the dependent variable and the array of values for the independent variable.

With the model created, we calculate the regression coefficients using the model's [fit\(\)](http://www.statsmodels.org/dev/generated/statsmodels.regression.linear_model.OLS.fit.html#statsmodels.regression.linear_model.OLS.fit) (http://www.statsmodels.org/dev/generated/statsmodels.regression.linear_model.OLS.fit.html#statsmodels.regression.linear_model.OLS.fit) method; this method returns a [RegressionResults](http://www.statsmodels.org/dev/generated/statsmodels.regression.linear_model.RegressionResults.html#statsmodels.regression.linear_model.RegressionResults) (http://www.statsmodels.org/dev/generated/statsmodels.regression.linear_model.RegressionResults.html#statsmodels.regression.linear_model.RegressionResults) object. To display information about the fit, we display the output of the result's [summary\(\)](http://www.statsmodels.org/dev/generated/statsmodels.regression.linear_model.RegressionResults.summary.html#statsmodels.regression.linear_model.RegressionResults.summary) (http://www.statsmodels.org/dev/generated/statsmodels.regression.linear_model.RegressionResults.summary.html#statsmodels.regression.linear_model.RegressionResults.summary) method.

```
In [109]: model = sm.OLS(y, X)
          results = model.fit()
          display(results.summary())
```

OLS Regression Results

Dep. Variable:	y	R-squared:	0.995
Model:	OLS	Adj. R-squared:	0.995
Method:	Least Squares	F-statistic:	1.906e+04
Date:	Tue, 17 Apr 2018	Prob (F-statistic):	4.40e-114
Time:	13:37:57	Log-Likelihood:	-145.36
No. Observations:	100	AIC:	294.7
Df Residuals:	98	BIC:	299.9
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	3.1799	0.208	15.317	0.000	2.768	3.592
x1	4.9515	0.036	138.049	0.000	4.880	5.023

Omnibus:	7.365	Durbin-Watson:	1.662
Prob(Omnibus):	0.025	Jarque-Bera (JB):	6.878
Skew:	0.569	Prob(JB):	0.0321
Kurtosis:	3.596	Cond. No.	11.7

From the results, the coefficients, their p-values, and the value of `R-squared` are particularly of interest. We can access these directly from `results` using the `params`, `pvalues`, and `rsquared` properties as well.

```
In [110]: print('Parameters: ', results.params)
          print("P-values:", results.pvalues)
          print('R^2: ', results.rsquared)

Parameters: [3.17992448 4.95145343]
P-values: [9.41199538e-028 4.40483504e-114]
R^2: 0.9948839945191904
```

The coefficients appear in the same order in which the independent variables appear in `x`.


Note that this is "close" to the equation we used to generate the data - the discrepancy is due to the error we introduced.

Let's plot the regression line along with our data. We can repeat the same steps as before to create the scatter plot. We make an additional call to `plot()` and specify the y-values as the output from the `results.predict()` method which returns the predicted values of the dependent variable based on the

regression results. Specifying 'r' for the third argument to plot indicates that we would like the line to be red.

```
In [111]: fig, ax = plt.subplots(figsize=(12,10))
          ax.plot(x, y, 'b+')
          ax.plot(x, results.predict(), 'r')
```

```
Out[111]: [<matplotlib.lines.Line2D at 0x10feb6b70>]
```



The [p-value](https://en.wikipedia.org/wiki/P-value) (<https://en.wikipedia.org/wiki/P-value>) associated with each coefficient indicates how likely changes to the corresponding independent variable account for changes in the dependent variable. A p-value close to zero (typically, less than 0.05) indicates that the independent variable provides a meaningful addition to the model. ¹

The R-squared value is also known as the [coefficient of determination](https://en.wikipedia.org/wiki/Coefficient_of_determination) (https://en.wikipedia.org/wiki/Coefficient_of_determination) and provides a measure of how well the regression line fits the data. The coefficient of determination can range from 0 to 1 with 0 indicating (in some regressions, the value can be negative) and indicates how much of the variation in dependent variable can be explained by the model - a value of 1 indicates that all variation is explained by the model and 0 indicates that the model accounts for none of the variation. In this example, some of the variation is due to the error we introduced - changing the magnitude of the error or the parameters of the distribution from which values were drawn will result in a better or worse coefficient of determination.

Assuming the fit is a good one, we can use it predict other values not in the initial data. The simplest way to do this for individual values is by defining a function that takes an input value as a parameter and computes and returns the output based on the calculate coefficients.

```
In [112]: def predict(model_results, value):  
          const, coeff = model_results.params  
          return const + coeff * value  
  
print(predict(results, 4.5))  
print(predict(results, 15))
```

```
25.461464939427398  
77.45172600059679
```

Simple Linear Regression

Let's look at an example based on real data. To begin, we'll load the fuel economy data we processed previously.

```
In [113]: epa_data = pd.read_csv("./data/02-vehicles.csv", engine="python")
```

Recall that a description of the data is available in `./data/02-vehicles-description.html`.

```
In [114]: from IPython.display import HTML  
HTML(filename="./data/02-vehicles-description.html")
```

- atvtype - type of alternative fuel or advanced technology vehicle
 - Bifuel (CNG) - Bi-fuel gasoline and compressed natural gas vehicle
 - Bifuel (LPG) - Bi-fuel gasoline and propane vehicle
 - CNG - Compressed natural gas vehicle
 - Diesel - Diesel vehicle
 - EV - Electric vehicle
 - FFV - Flexible fueled vehicle (gasoline or E85)
 - Hybrid - Hybrid vehicle
 - Plug-in Hybrid - Plug-in hybrid vehicle
- barrels08 - annual petroleum consumption in barrels for fuelType1
- barrelsA08 - annual petroleum consumption in barrels for fuelType2
- charge120 - time to charge an electric vehicle in hours at 120 V
- charge240 - time to charge an electric vehicle in hours at 240 V
- city08 - city MPG for fuelType1
- city08U - unrounded city MPG for fuelType1
- cityA08 - city MPG for fuelType2

For this unit, we'll work the the following columns.

- co2
- comb08
- cylinders

- displ
- highway08
- city08

We can also remove rows with missing or non-positive values. Because the DataFrame consists of only numeric data, we can use a sort of shortcut for removing rows with non-positive values. We create a mask applied to the entire DataFrame rather than a specific column - this applies it to all columns. We then use `all()` with an argument of 1 to indicate that the property applies across all columns. Effectively this mask will match any row in which all the values are positive.

```
In [115]: epa_subset = epa_data[['co2', 'comb08', 'cylinders', 'displ', 'highway08', 'city08']]
epa_subset = epa_subset[(epa_subset > 0).all(1)].copy()
epa_subset.head()
```

Out[115]:

	co2	comb08	cylinders	displ	highway08	city08
16780	318	32	4.0	2.0	37	29
16781	315	32	4.0	2.0	39	29
16839	318	32	4.0	2.0	37	29
16840	315	32	4.0	2.0	39	29
21337	315	32	4.0	2.0	39	29

To start, let's look at the `co2`, `comb08`, `cylinders`, and `displ` columns in `epa_subset`. First we can calculate the correlation coefficient matrix. With a DataFrame, we can use the `corr()` method to calculate this.

```
In [116]: epa_subset[['co2', 'comb08', 'cylinders', 'displ']].corr()
```

Out[116]:


	co2	comb08	cylinders	displ
co2	1.000000	-0.928998	0.824112	0.850303
comb08	-0.928998	1.000000	-0.743347	-0.779056
cylinders	0.824112	-0.743347	1.000000	0.927088
displ	0.850303	-0.779056	0.927088	1.000000

To get sense of these relationships, visually, we can use a pair plot.

Lab 1 In the cell below, create a pair plot for the `co2`, `comb08`, `cylinders`, and `displ` columns in `epa_subset`.

```
In [117]: sns.pairplot(data=epa_subset[['co2', 'comb08', 'cylinders', 'displ']])
```

```
Out[117]: <seaborn.axisgrid.PairGrid at 0x1120b6f98>
```




Let's look more closely at the relationship between `cylinders` and `displacement`.

Lab 2 In the cell below, create scatter plot for the `cylinders` and `displacement` columns in `epa_subset`.

```
In [118]: epa_subset.plot.scatter(x='cylinders', y='displ')
```

```
Out[118]: <matplotlib.axes._subplots.AxesSubplot at 0x104052f28>
```



Next, lets create the least squares model and fit a line to the data.

```
In [119]: X = sm.add_constant(epa_subset.cylinders)
Y = epa_subset.displ
model = sm.OLS(Y, X)
res = model.fit()
display(res.summary())
```

OLS Regression Results

Dep. Variable:	displ	R-squared:	0.859
Model:	OLS	Adj. R-squared:	0.859
Method:	Least Squares	F-statistic:	4.527e+04
Date:	Tue, 17 Apr 2018	Prob (F-statistic):	0.00
Time:	13:38:02	Log-Likelihood:	-5593.4
No. Observations:	7402	AIC:	1.119e+04
Df Residuals:	7400	BIC:	1.120e+04
Df Model:	1		
Covariance Type:	nonrobust		


	coef	std err	t	P> t	[0.025	0.975]
const	-0.6751	0.019	-35.240	0.000	-0.713	-0.638
cylinders	0.6834	0.003	212.759	0.000	0.677	0.690

Omnibus:	488.472	Durbin-Watson:	0.995
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1327.138
Skew:	0.364	Prob(JB):	6.54e-289
Kurtosis:	4.942	Cond. No.	19.6

From the `R-squared` value we can see there is a somewhat strong linear relationship between the data; further, the p-value for the coefficient of `cylinders` indicates that changes in `displ` are likely attributed to changes in `cylinders`.

To plot the fit line with the scatter plot generated by the DataFrame's `plot()` method we can use StatsModel's [`abline_plot\(\)`](http://www.statsmodels.org/dev/generated/statsmodels.graphics.regressionplots.abline_plot.html) function. First, we import the function then create a scatter plot and store the returned axes object.


```
In [120]: from statsmodels.graphics.regressionplots import abline_plot
axes = epa_subset.plot.scatter(x="cylinders", y="displ")
```



We can add the plot of the regression line to the plot using `abline_plot()` function, specifying the model results, the axes, and a color.

```
In [121]: abline_plot(model_results=res, ax=axes, color='r')
```

Out[121]:




Often when creating models, we want to exclude outliers in our calculations. Let's look at the box plots for `cylinders` and `displ`.

Lab 3 In the cells below, create the box plots for `cylinders` and `displ`.


```
In [122]: epa_subset.cylinders.plot(kind='box')
```

```
Out[122]: <matplotlib.axes._subplots.AxesSubplot at 0x114ddeb70>
```




```
In [123]: epa_subset.displ.plot(kind='box')
```

```
Out[123]: <matplotlib.axes._subplots.AxesSubplot at 0x1130d83c8>
```



We can see that there are a few outliers for each column. We can remove them and recalculate the fit. We can create a copy of the DataFrame from which we will remove outliers.

```
In [124]: epa_no_outliers = epa_subset[['cylinders', 'displ']].copy()
```

To remove the outliers from a column, we'll create a function that we can apply to a DataFrame. We'll define outliers based on the interquartile range.

After defining the function, we can use it with our `epa_no_outliers` DataFrame.

```
In [125]: def remove_outliers(dataframe, column):
    q1 = dataframe[column].quantile(0.25)
    q3 = dataframe[column].quantile(0.75)
    iqr = q3 - q1
    lower = q1 - 1.5 * iqr
    upper = q3 + 1.5 * iqr
    return dataframe[(dataframe[column] >= lower) &
                     (dataframe[column] <= upper)].copy()

epa_no_outliers = remove_outliers(epa_no_outliers, "cylinders")
epa_no_outliers = remove_outliers(epa_no_outliers, "displ")
```

With the outliers removed, we can create a new model and calculate the regression coefficients.

Lab 4 In the cell below, create an ordinary least squares model where `cylinders` is the independent variable and `displ` is the dependent variable. Include a coefficient term in the model. Calculate the fit and store the result in a variable named `res_no_outliers`.

```
In [126]: X = sm.add_constant(epa_no_outliers.cylinders)
Y = epa_no_outliers.displ
model = sm.OLS(Y, X)
res_no_outliers = model.fit()
```

Looking at the result's summary and the `R-squared` value, we can see the the fit is slightly better with the outliers removed.

```
In [127]: display(res_no_outliers.summary())
```

OLS Regression Results


Dep. Variable:	displ	R-squared:	0.883
Model:	OLS	Adj. R-squared:	0.883
Method:	Least Squares	F-statistic:	5.384e+04
Date:	Tue, 17 Apr 2018	Prob (F-statistic):	0.00
Time:	13:38:03	Log-Likelihood:	-4103.9
No. Observations:	7132	AIC:	8212.
Df Residuals:	7130	BIC:	8225.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-1.1141	0.019	-59.314	0.000	-1.151	-1.077
cylinders	0.7681	0.003	232.042	0.000	0.762	0.775

Omnibus:	104.759	Durbin-Watson:	1.014
Prob(Omnibus):	0.000	Jarque-Bera (JB):	153.838
Skew:	0.167	Prob(JB):	3.93e-34
Kurtosis:	3.637	Cond. No.	21.5


We can plot both the original fit and the fit calculated without outliers against a scatter plot of the data.

```
In [128]: axes = epa_subset.plot.scatter(x="cylinders", y="displ")
```



```
In [129]: abline_plot(model_results=res, ax=axes, color='r')
          abline_plot(model_results=res_no_outliers, ax=axes, color='b')
```

Out[129]:





The original regression line is colored red and the regression line calculated with outliers removed is colored blue.

Multiple Linear Regression

So far, we've looked at regressions in which there is one independent variable and a constant. Often, changes in the response variable are dependent on multiple variables. For example, we expect that the combined fuel economy is dependent on both city and highway economy. We can see from the scatter plots that `comb08` looks linearly dependent on `city08` and `highway08`.

```
In [130]: epa_subset.plot.scatter(x='city08', y='comb08')
epa_subset.plot.scatter(x='highway08', y='comb08')
```

```
Out[130]: <matplotlib.axes._subplots.AxesSubplot at 0x113a11eb8>
```



We can also fit models for each of these individually - `city08 / comb08` and

highway08 / comb08 .

```
In [131]: X = sm.add_constant(epa_subset.city08)
Y = epa_subset.comb08
model = sm.OLS(Y, X)
res = model.fit()
display(res.summary())
```

OLS Regression Results

Dep. Variable:	comb08	R-squared:	0.972
Model:	OLS	Adj. R-squared:	0.972
Method:	Least Squares	F-statistic:	2.590e+05
Date:	Tue, 17 Apr 2018	Prob (F-statistic):	0.00
Time:	13:38:04	Log-Likelihood:	-10242.
No. Observations:	7402	AIC:	2.049e+04
Df Residuals:	7400	BIC:	2.050e+04
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	2.9449	0.041	71.791	0.000	2.865	3.025
city08	0.9854	0.002	508.942	0.000	0.982	0.989

Omnibus:	1685.615	Durbin-Watson:	1.292
Prob(Omnibus):	0.000	Jarque-Bera (JB):	6373.261
Skew:	-1.096	Prob(JB):	0.00
Kurtosis:	6.982	Cond. No.	77.6

```
In [132]: X = sm.add_constant(epa_subset.highway08)
Y = epa_subset.comb08
model = sm.OLS(Y, X)
res = model.fit()
display(res.summary())
```

OLS Regression Results

Dep. Variable:	comb08	R-squared:	0.925
Model:	OLS	Adj. R-squared:	0.925
Method:	Least Squares	F-statistic:	9.094e+04
Date:	Tue, 17 Apr 2018	Prob (F-statistic):	0.00
Time:	13:38:04	Log-Likelihood:	-13930.
No. Observations:	7402	AIC:	2.786e+04
Df Residuals:	7400	BIC:	2.788e+04
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-2.6589	0.087	-30.508	0.000	-2.830	-2.488
highway08	0.9307	0.003	301.560	0.000	0.925	0.937

Omnibus:	2766.126	Durbin-Watson:	1.294
Prob(Omnibus):	0.000	Jarque-Bera (JB):	17613.587
Skew:	1.653	Prob(JB):	0.00
Kurtosis:	9.796	Cond. No.	133.

Separately, the models fit the data quite well. Let's look at the model in which both `city08` and `highway08` are independent variables.

```
In [133]: X = sm.add_constant(epa_subset[['city08', 'highway08']])
Y = epa_subset.comb08
model = sm.OLS(Y, X)
res = model.fit()
display(res.summary())
```

OLS Regression Results

Dep. Variable:	comb08	R-squared:	0.996
Model:	OLS	Adj. R-squared:	0.996
Method:	Least Squares	F-statistic:	8.489e+05
Date:	Tue, 17 Apr 2018	Prob (F-statistic):	0.00
Time:	13:38:04	Log-Likelihood:	-3370.7
No. Observations:	7402	AIC:	6747.
Df Residuals:	7399	BIC:	6768.
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-0.0860	0.022	-3.874	0.000	-0.130	-0.042
city08	0.6466	0.002	347.731	0.000	0.643	0.650
highway08	0.3600	0.002	199.909	0.000	0.356	0.364

Omnibus:	101.613	Durbin-Watson:	1.802
Prob(Omnibus):	0.000	Jarque-Bera (JB):	62.826
Skew:	-0.059	Prob(JB):	2.28e-14
Kurtosis:	2.564	Cond. No.	177.

Looking at the coefficient of determination, we can see this model, which depends on both `city08` and `highway08`, fits the data better than a model which depends on only one of the variables.

We can predict values using the results in a method similar to before using a custom function that accepts values for each of the independent variables.

```
In [134]: def predict(model_results, city, highway):
const, city_coeff, highway_coeff = model_results.params
return const + city_coeff * city + highway_coeff * highway

print(predict(res, 25, 35))
print(predict(res, 45, 50))
```

```
28.67866070951573
47.01052930817669
```


For our problem, the results indicate that if a vehicle's city economy is 25 mpg and the highway economy is 35 mpg, the combined economy is about 29 mpg. Similarly, if the city and highway economies are 45 mpg and 50 mpg, respectively, the combined economy is about 47 mpg.


Linear-Like Regression

While there are many of relationships that are linear and that can be modeled using a linear regression, there are also relationships that are non-linear. Among these non-linear relationships are those that can be transformed into a linear ones in terms of the coefficients and independent variables.

As an example, consider the relationship between `comb08` and `co2`.

```
In [135]: epa_subset.plot.scatter(x="comb08", y="co2")
```

```
Out[135]: <matplotlib.axes._subplots.AxesSubplot at 0x10e34fcf8>
```



While this relationship doesn't appear to be linear, it does look [hyperbolic](https://en.wikipedia.org/wiki/Hyperbola) (<https://en.wikipedia.org/wiki/Hyperbola>). In this case, the relationship between the independent variable and dependent variable could be written as

$$Y = \frac{1}{\beta_0 + \beta_1 X}$$

To move the coefficients and independent variable out of the denominator, we can take the reciprocal of both sides (provided the neither side is zero) - this gives us

$$\frac{1}{Y} = \beta_0 + \beta_1 X$$

For a given observation, this form is easier to work with since $\frac{1}{Y}$ and X are constants and we need to solve for β_0 and β_1 .

We can calculate the reciprocal of the dependent variable, `co2`, and store the value in a new column.

```
In [136]: epa_non_linear = epa_subset[['comb08', 'co2']].copy()
```

To transform the problem into a linear one, we need to calculate the reciprocal of the `co2` values.


```
In [137]: epa_non_linear["reciprocal_co2"] = 1/epa_non_linear.co2
```

Looking at the plot of `comb08` and `reciprocal_co2`, it appears that the relationship is linear, which supports our assumption that the original relationship was hyperbolic.

Lab 5 In the cell below, create a scatter plot of `comb08` and `reciprocal_co2`.

```
In [138]: epa_non_linear.plot.scatter(x="comb08", y="reciprocal_co2")
```

```
Out[138]: <matplotlib.axes._subplots.AxesSubplot at 0x113a1e080>
```




Before fitting the data with a linear model, let's remove the outliers.

Lab 6 In the cell below, remove the outliers from the `comb08` and `reciprocal_co2` columns in the `epa_non_linear` DataFrame. Use the `remove_outliers()` function we created earlier.

```
In [139]: epa_non_linear = remove_outliers(epa_non_linear, "reciprocal_co2")
```

With the outliers removed, let's look at the scatter plot again.

```
In [140]: axes = epa_non_linear.plot.scatter(x="comb08", y="reciprocal_co2")
```



We can now create a linear model for `comb08` and `reciprocal_co2`. After calculating the coefficients, we display the summary of the results.

```
In [141]: X = sm.add_constant(epa_non_linear["comb08"])
Y = epa_non_linear["reciprocal_co2"]
model = sm.OLS(Y, X)
res = model.fit()
display(res.summary())
```

OLS Regression Results


Dep. Variable:	reciprocal_co2	R-squared:	0.974
Model:	OLS	Adj. R-squared:	0.974
Method:	Least Squares	F-statistic:	2.656e+05
Date:	Tue, 17 Apr 2018	Prob (F-statistic):	0.00
Time:	13:38:05	Log-Likelihood:	56904.
No. Observations:	7218	AIC:	-1.138e+05
Df Residuals:	7216	BIC:	-1.138e+05
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	6.93e-05	4.91e-06	14.115	0.000	5.97e-05	7.89e-05
comb08	0.0001	2.12e-07	515.335	0.000	0.000	0.000

Omnibus:	5458.351	Durbin-Watson:	1.533
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1617181.198
Skew:	2.579	Prob(JB):	0.00
Kurtosis:	76.148	Cond. No.	106.

From the coefficient of determination, we see that the model produced a good fit. Adding the regression line to the existing scatter plot give the following plot.

```
In [142]: abline_plot(model_results=res, ax=axes, color='r')
```



We now need to transform the regression line to fit the original data. Using the *params* property of the results we have the following constant term and coefficient.

```
In [143]: res.params
```

```
Out[143]: const      0.000069  
comb08      0.000109  
dtype: float64
```

This means that our model is

$$Y = \frac{1}{0.000068 + 0.000109X}$$


We can calculate the predicted values from our model using the following code. We use *sort_values()* to ensure that the values of the independent variable are in order. This will be important when we plot the curve; we didn't need to do this previously as we relied on the *abline_plot()* function, which handled this for us,

```
In [144]: prediction = 1/(res.params.const +  
                           res.params.comb08 * epa_subset.comb08.sort_values())
```

We can now plot the model curve against our original data.

```
In [145]: fig, axes = plt.subplots(figsize=(12,10))
          axes.plot(epa_subset.comb08, epa_subset.co2, 'b+')
          axes.plot(epa_subset.comb08.sort_values(), prediction, 'r')
```

```
Out[145]: [<matplotlib.lines.Line2D at 0x10f2d42b0>]
```



Logistic Regression

A [logistic regression](https://en.wikipedia.org/wiki/Logistic_regression) (https://en.wikipedia.org/wiki/Logistic_regression) is used to model data where the dependent variable is categorical. In the simplest case, the dependent variable is binary and has only two possible values. A logistic model, provides an estimate of the probability that one of the two categories applies given the values of the independent variables; it fits a [logistic probability distribution](https://en.wikipedia.org/wiki/Logistic_distribution) (https://en.wikipedia.org/wiki/Logistic_distribution) to the data. We'll only look at simple case where the dependent variable is binary and there is only one independent variable.

An Example

As an example, consider the the following [example taken from the Wikipedia page on logistic regressions](https://en.wikipedia.org/wiki/Logistic_regression#Example:_Probability_of_passing_an_exam_versus_hours) ([https://en.wikipedia.org/wiki/Logistic_regression#Example: Probability of passing an exam versus](https://en.wikipedia.org/wiki/Logistic_regression#Example:_Probability_of_passing_an_exam_versus_hours) [https://en.wikipedia.org/wiki/Logistic_regression#Example: Probability of passing an exam versus](https://en.wikipedia.org/wiki/Logistic_regression#Example:_Probability_of_passing_an_exam_versus_hours)). We have two variables: *hours* and *passed*. The *hours* variable represents the number of hours a


student spent studying for an exam and *passed* indicates whether or not the student passed the exam where 0 indicates failure and 1 indicates that the student passed; *hours* is a continuous variable and *passed* is a discrete, binary variable.

```
In [146]: hours = [0.5, 0.75, 1.0, 1.25, 1.50, 1.75, 1.75, 2.0, 2.25, 2.5, 2.75, 3.0,
passed = [0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1]
```

Plotting this data as a scatter plot give the following.

```
In [147]: plt.plot(hours, passed, 'b+')
```

```
Out[147]: [<matplotlib.lines.Line2D at 0x10e2c2cc0>]
```



The logistic regression calculates values for β_0 and β_1 in the following formula

$$P = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X)}}$$

where P is a probability value between 0 and 1 and X is the independent variable.

We can use the StatsModels [Logit\(\)](#)

(http://www.statsmodels.org/dev/generated/statsmodels.discrete.discrete_model.Logit.html)

function to perform the logistic regression.

We have to reassign the value of `stats.chisqprob` due to a discrepancy between the StatsModels module and the libraries on which it depends.

```
In [148]: from scipy import stats
stats.chisqprob = lambda chisq, df: stats.chi2.sf(chisq, df)

X = sm.add_constant(hours)
logit_model=sm.Logit(passed,X)
result=logit_model.fit()
display(result.summary())
```

```
Optimization terminated successfully.
      Current function value: 0.401494
      Iterations 7
```

Logit Regression Results

Dep. Variable:	y	No. Observations:	20
Model:	Logit	Df Residuals:	18
Method:	MLE	Df Model:	1
Date:	Tue, 17 Apr 2018	Pseudo R-squ.:	0.4208
Time:	13:38:06	Log-Likelihood:	-8.0299
converged:	True	LL-Null:	-13.863
		LLR p-value:	0.0006365

	coef	std err	z	P> z	[0.025	0.975]
const	-4.0777	1.761	-2.316	0.021	-7.529	-0.626
x1	1.5046	0.629	2.393	0.017	0.272	2.737


We create the model in much the same way as we did for a linear model. We specify the independent variable and add a constant using StatsModels' `add_constant()` function. We next create a logistic model using the `Logit()` function. To calculate the coefficients, we use model's `fit()` method. After the calculation is complete, we can view a summary of the results.

Just like the linear model, the results of the logistic model have a `predict()` method that give the model's predicted values based on the values of the independent variable. We can use this to plot the logistic curve along with the scatter plot of the data.

The model represents the probability of passing the exam given some number of hours spent studying.


```
In [149]: fig, axes = plt.subplots(figsize=(10,8))
axes.plot(hours,passed, 'b+')
axes.plot(hours, result.predict(), 'r-')
```

```
Out[149]: [<matplotlib.lines.Line2D at 0x112804898>]
```



We can calculate the probability of for a given value of the independent variable using a function.

```
In [150]: def probability(model_results, value):
            const, coeff = model_results.params
            exponent = -(const + coeff * value)
            denominator = 1 + np.exp(exponent)
            return 1 / denominator

            print(probability(result, 2))
            print(probability(result, 6))
```

```
0.25570318264090985
0.9929675242040855
```

For this problem, the results indicate if a student spends 2 hours studying there is about a 26% likelihood that the student will pass the exam; similarly if the student spends 6 hours studying, the probability of passing is greater than 99%.

Home Data

For an example with real data, consider the count auditor data we worked with previously. We can load the data from our local database.

```
In [151]: from sqlalchemy import create_engine
engine = create_engine('sqlite:///data/output.sqlite')
home_data = pd.read_sql("home_data", con=engine)
home_data.head()
```

Out[151]:

	index	AirConditioning	AppraisedBuilding	AppraisedLand	Area	Bathrooms	Bedrooms	County
0	0	True	59600.0	8100.0	2264	2.0	4.0	Franklin
1	1	True	69800.0	4600.0	1835	1.5	4.0	Franklin
2	2	True	60600.0	4900.0	1656	1.0	3.0	Franklin
3	3	True	31200.0	5000.0	1000	1.0	2.0	Franklin
4	4	True	63300.0	4600.0	1306	2.0	4.0	Franklin

For this example, lets see if we can calculate the logistic model that gives the probability of a house having a fireplace given its area. Currently, the dataset includes the the number of fireplaces a given property has; we need to convert values greater than zero to 1, indicating that there is a fireplace.

First, we drop any rows with missing data. Next, we create a new column, `HasFireplace` that is equal to the mask corresponding the `Fireplaces` being greater than zero. Masks return values of `True` or `False` and the `astype(int)` function call will convert the boolean value to an integer where `False` becomes 0 and `True` becomes 1.

```
In [152]: home_data.dropna(inplace=True)
home_data['HasFireplace'] = (home_data.Fireplaces > 0).astype(int)
home_data.head()
```

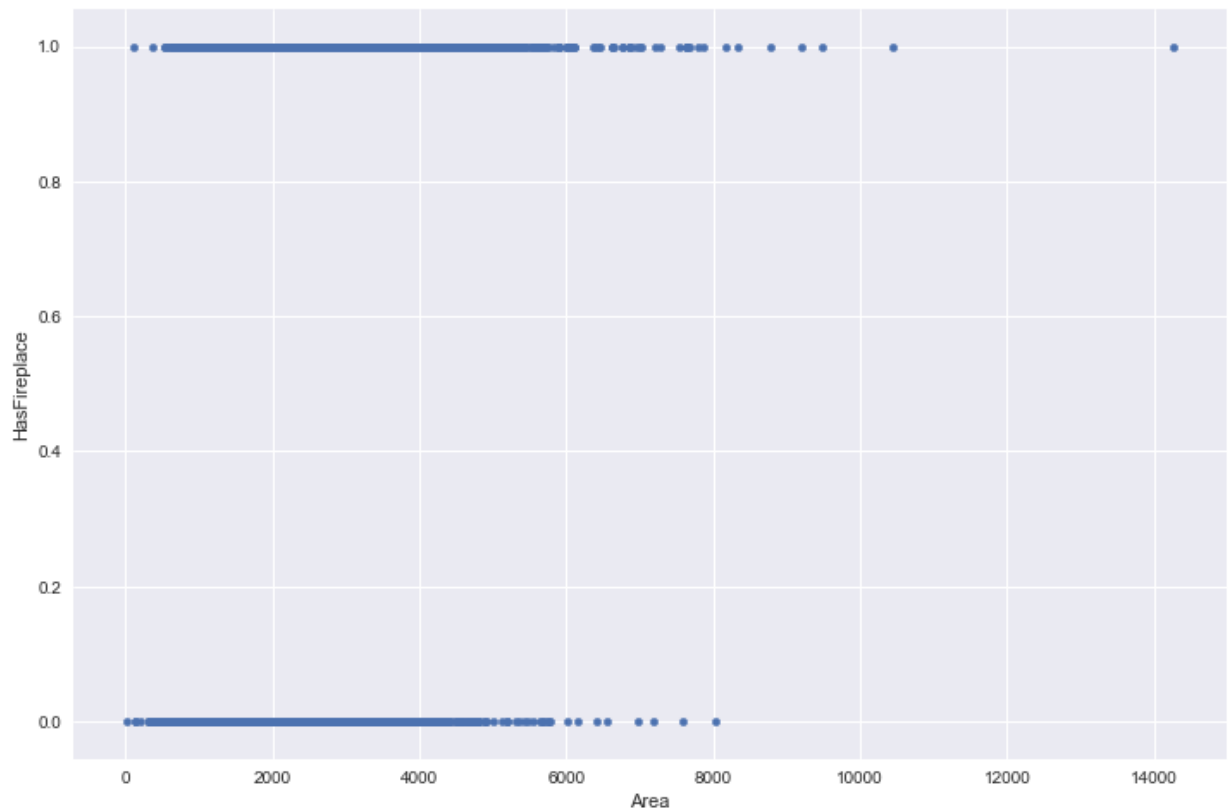
Out[152]:

	index	AirConditioning	AppraisedBuilding	AppraisedLand	Area	Bathrooms	Bedrooms	County
0	0	True	59600.0	8100.0	2264	2.0	4.0	Franklin
1	1	True	69800.0	4600.0	1835	1.5	4.0	Franklin
2	2	True	60600.0	4900.0	1656	1.0	3.0	Franklin
3	3	True	31200.0	5000.0	1000	1.0	2.0	Franklin
4	4	True	63300.0	4600.0	1306	2.0	4.0	Franklin

We can now create the scatter plot of `Area` and `HasFireplaces`

```
In [153]: home_data.plot.scatter(x="Area", y="HasFireplace")
```

```
Out[153]: <matplotlib.axes._subplots.AxesSubplot at 0x110c52ba8>
```



Before creating the model, we sort the values of our independent variable; this will aid in plotting later.

```
In [154]: home_data.sort_values(by=["Area"], inplace=True)
```

Creating the logistic model and calculating the regression coefficients is similar to the logistic model as we noted in the example.

```
In [155]: X = sm.add_constant(home_data.Area)
logit_model=sm.Logit(home_data.HasFireplace,X)
result=logit_model.fit()
display(result.summary())
```

```
Optimization terminated successfully.
      Current function value: 0.636377
      Iterations 5
```

Logit Regression Results

Dep. Variable:	HasFireplace	No. Observations:	39236
Model:	Logit	Df Residuals:	39234
Method:	MLE	Df Model:	1
Date:	Tue, 17 Apr 2018	Pseudo R-squ.:	0.05027
Time:	13:38:08	Log-Likelihood:	-24969.
converged:	True	LL-Null:	-26291.
		LLR p-value:	0.000

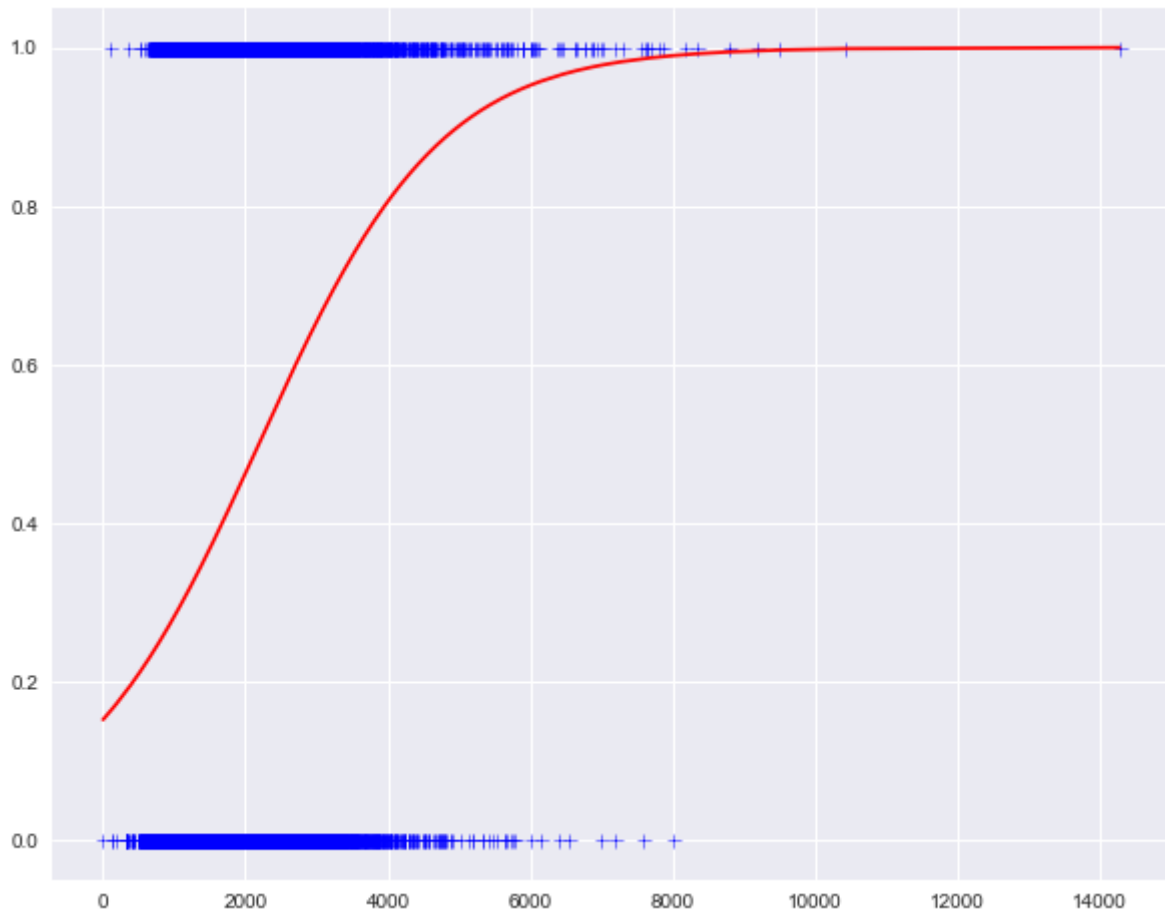
	coef	std err	z	P> z	[0.025	0.975]
const	-1.7215	0.029	-59.697	0.000	-1.778	-1.665
Area	0.0008	1.64e-05	47.992	0.000	0.001	0.001

With the model created and the coefficients calculated, we can now plot the regression curve against the original data.

Lab 7 In the cell below, create a scatter plot of the `Area` and `HasFireplace` columns from the `home_data` DataFrame. On the same figure, plot the logistic regression curve.

```
In [156]: fig, axes = plt.subplots(figsize=(10,8))
axes.plot(home_data.Area, home_data.HasFireplace, 'b+')
axes.plot(home_data.Area, result.predict(), 'r-')
```

```
Out[156]: [<matplotlib.lines.Line2D at 0x10e744dd8>]
```



The formula for the curve is given by

$$P = \frac{1}{1+e^{-(-1.7215+0.0008X)}}$$

For a given area, we can calculate the probability that the house has a fireplace using this formula based on the model.

Let's look at one more example. Older home tend to have fewer bathrooms. We can create a new column, `MoreThanOneBathroom`

Lab 8 In the cell below, create a new column named `MoreThanOneBathroom` in the `home_data` DataFrame that has a value of 0 if the value of `Bathrooms` is less than or equal to 1 and has a value of 1 otherwise.

```
In [157]: home_data["MoreThanOneBathroom"] = (home_data.Bathrooms > 1).astype(int)
```

We should sort the data by `YearBuilt` before continuing.

```
In [158]: home_data.sort_values(by=["YearBuilt"], inplace=True)
```

As before we construct the model with `YearBuilt` as the independent variable and `MoreThanOneBathroom` as the dependent variable. We calculate a logistic curve that best fits the data.

```
In [159]: X = sm.add_constant(home_data.YearBuilt)
logit_model=sm.Logit(home_data.MoreThanOneBathroom, X)
result=logit_model.fit()
display(result.summary())
```

```
Optimization terminated successfully.
      Current function value: 0.490881
      Iterations 6
```

Logit Regression Results

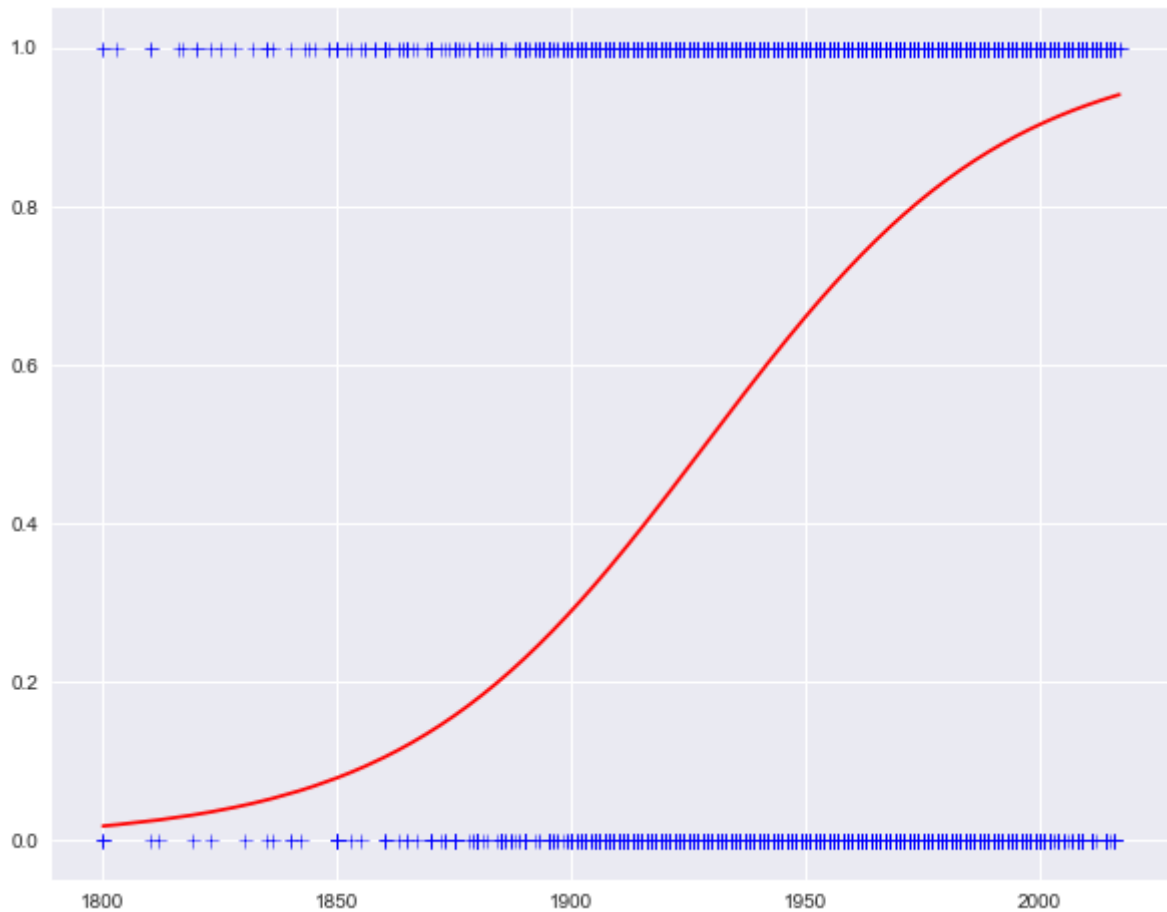
Dep. Variable:	MoreThanOneBathroom	No. Observations:	39236
Model:	Logit	Df Residuals:	39234
Method:	MLE	Df Model:	1
Date:	Tue, 17 Apr 2018	Pseudo R-squ.:	0.1333
Time:	13:38:08	Log-Likelihood:	-19260.
converged:	True	LL-Null:	-22222.
		LLR p-value:	0.000

	coef	std err	z	P> z	[0.025	0.975]
const	-60.3612	0.872	-69.209	0.000	-62.071	-58.652
YearBuilt	0.0313	0.000	70.319	0.000	0.030	0.032

Finally, we can create a scatter plot of `YearBuilt` and `MoreThanOneBathroom` along with the logistic regression curve.

```
In [160]: fig, axes = plt.subplots(figsize=(10,8))
          axes.plot(home_data.YearBuilt, home_data.MoreThanOneBathroom, 'b+')
          axes.plot(home_data.YearBuilt, result.predict(), 'r-')
```

```
Out[160]: [<matplotlib.lines.Line2D at 0x114979cc0>]
```



Lab Answers

- `sns.pairplot(data=epa_subset[['co2', 'comb08', 'cylinders', 'displ']])`
 - `epa_subset.plot.scatter(x='cylinders', y='displ')`
 - `epa_subset.cylinders.plot(kind='box')`
- and
- ```
epa_subset.displ.plot(kind='box')
```
- ```
X = sm.add_constant(epa_no_outliers.cylinders)
Y = epa_no_outliers.displ
model = sm.OLS(Y, X)
res_no_outliers = model.fit()
```
 - `epa_non_linear.plot.scatter(x="comb08", y="reciprocal_co2")`

6.

```
epa_non_linear = remove_outliers(epa_non_linear, "comb08")
epa_non_linear = remove_outliers(epa_non_linear, "reciprocal_co
2")
```
7.

```
fig, axes = plt.subplots(figsize=(10,8))
axes.plot(home_data.Area, home_data.HasFireplace, 'b+')
axes.plot(home_data.Area, result.predict(), 'r-')
```
8.

```
home_data["MoreThanOneBathroom"] = (home_data.Bathrooms > 1).ast
ype(int)
```

Next Steps

The models we create can be used populate dashboards or included in reports. Later, we'll look at creating visualizations and reporting information. Model creation could also be an intermediate step in the analysis process; in a later unit we'll look at automating model creation.

Resources and Further Reading

- [Simple and Multiple Linear Regression in Python \(https://towardsdatascience.com/simple-and-multiple-linear-regression-in-python-c928425168f9\)](https://towardsdatascience.com/simple-and-multiple-linear-regression-in-python-c928425168f9)
- [Logistic Regression in Python Using Rodeo \(http://blog.yhat.com/posts/logistic-regression-python-rodeo.html\)](http://blog.yhat.com/posts/logistic-regression-python-rodeo.html)
- [Regression Analysis with Python by Massaron and Boschetti \(Safari Books\) \(http://proquest.safaribooksonline.com.csc.c.ohionet.org/book/programming/python/9781785286](http://proquest.safaribooksonline.com.csc.c.ohionet.org/book/programming/python/9781785286)
- [Data Science Algorithms in a Week by Natingga, Regression \(Safari Books\) \(http://proquest.safaribooksonline.com.csc.c.ohionet.org/book/programming/machine-learning/9781787284586/regression/1500cb6b_9703_4b4a_bffb_61da8fbd2e97_xhtml?unicode=ohlink\)](http://proquest.safaribooksonline.com.csc.c.ohionet.org/book/programming/machine-learning/9781787284586/regression/1500cb6b_9703_4b4a_bffb_61da8fbd2e97_xhtml?unicode=ohlink)

Notes

1. The null hypothesis being tested is that the variable associated with the coefficient has no effect on the dependent variable. When the p-value is sufficiently low, typically less than 0.05, we reject the null hypothesis thereby accepting that the variable does have an effect on the dependent variable.

Exercises

1. Calculate the coefficients for a linear model relating two columns from the fuel economy or county auditor data not discussed in the examples. Create a scatter plot of the data along with a plot of the regression line.

2. Calculate the coefficients for a logistic model relating two columns from the fuel economy or county auditor data not discussed in the examples. Create a scatter plot of the data along with a plot of the regression curve.

In []: